

MP4: The UNIX Shell

- `exec*()`
- `fork()`
- `wait()`
- `system()`
- `pipe`
- Background command
- Special command

exec*() family

- Basic

- `execl`(pathname, ...): take a list of args
- `execv`(pathname, argv): take an array of args
- E.g `int ret = execl("/bin/ls", "ls", "-l", (char*)0);`

- Custom environment

- `execle`(pathname, ..., envp)
- `execve`(pathname, argv, envp)
- E `char *env[] = { "HOME=/usr/home", "LOGNAME=home", (char*)0 };
int ret = execle("/bin/ls", "ls", "-l", (char*)0, env);`

- PATH

- `execvp`(filename, ...)
- `execvp`(filename, argv)
- E.g `char *cmd[] = { "ls", "-l", (char*)0 };
int ret = execvp("ls", cmd);`

exec*() family

- exec*() is a **brain transplant** function
 - Clears out the machine code of the calling program
 - Loads the code of the called program
 - Does not return on success

```
#include<stdio.h>

int main ()
{
    char * arglist [] = {"doesnotmatter", "-l", "-a"};
    printf ("<<<<<<< About to execute ls -l>>>>>>>\n");
    execvp ("ls", arglist);
    printf ("<<<<<<< ls is done >>>>>>>\n");
    return 0;
}
```

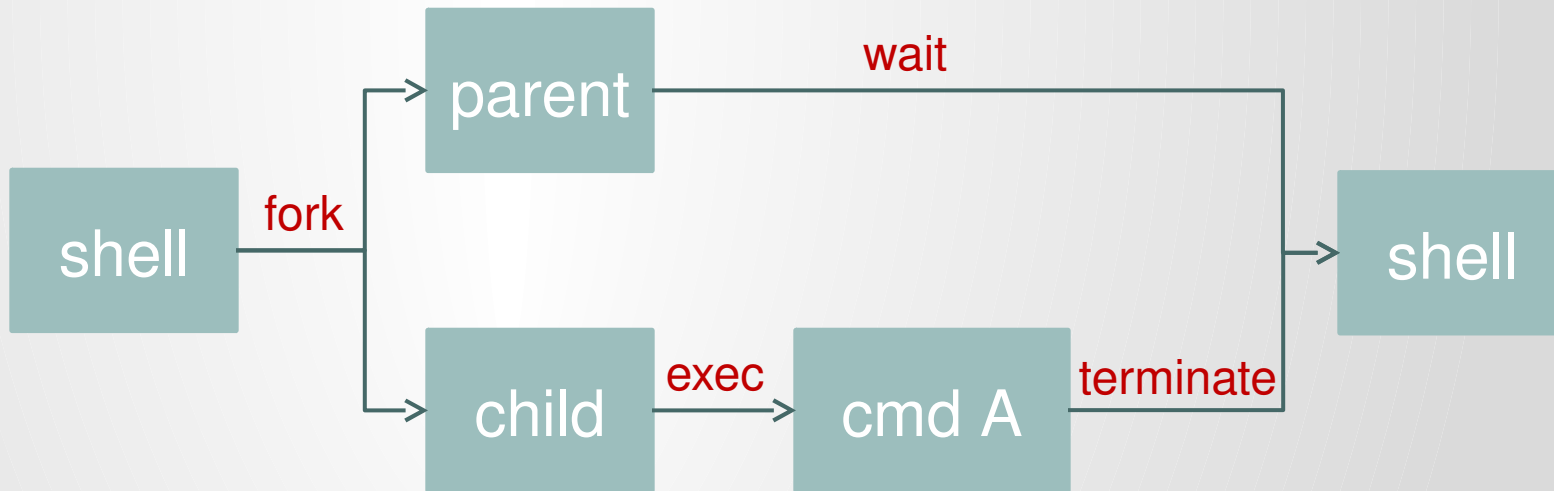
- Can we directly call exec() to handle a user's command in our shell program?
 - **NO! We'll lose our shell**
 - So what can we do?

fork()

- Create a new process
- Return twice if succeed
 - 0: child process
 - >0: parent process, the return number is the child's PID
- So we have a new process, then what?
 - Call `exec*()` in the **child** process
- Now the user's command is running, how can we wait until it finishes so that to accept next command?

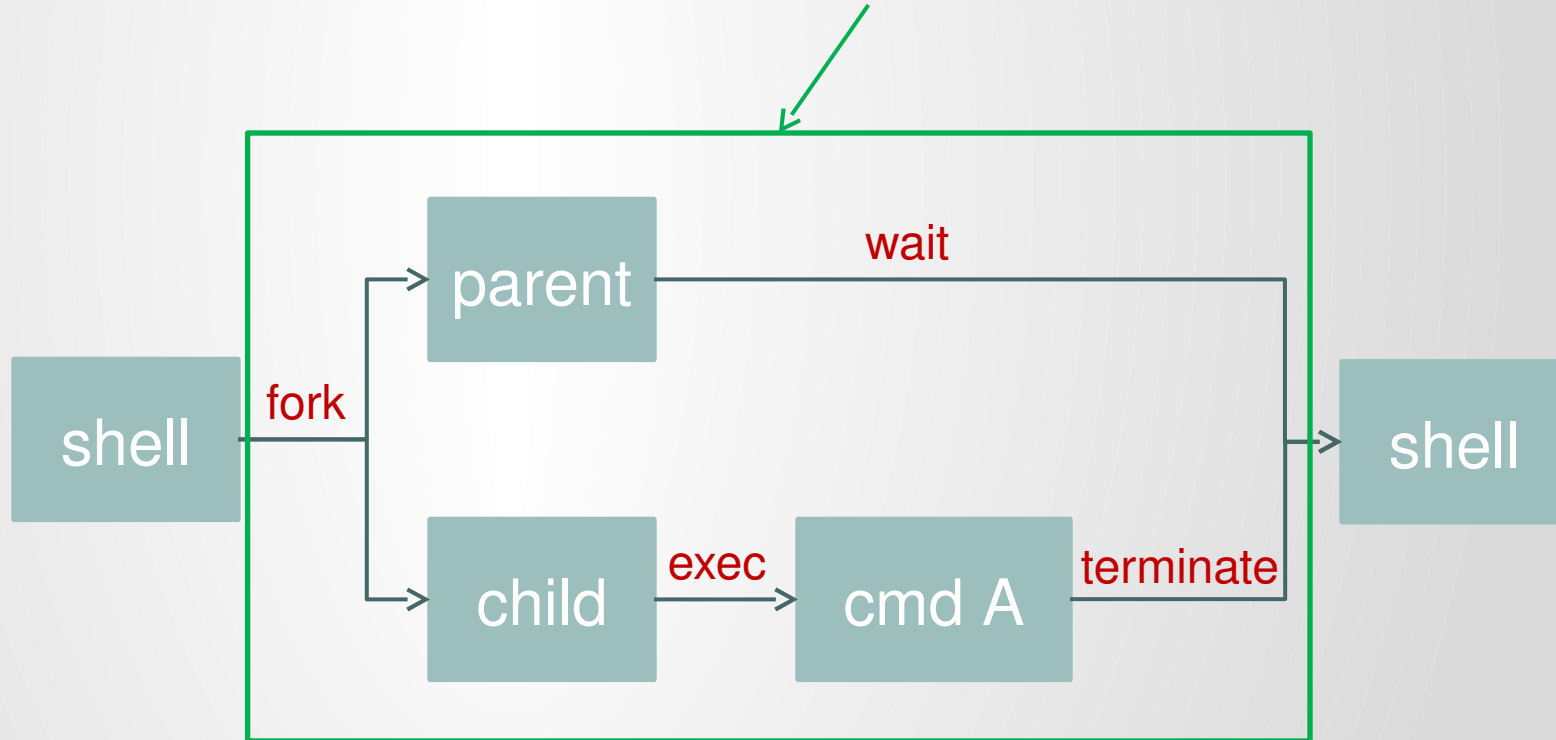
wait*()

- Synchronize with children
 - Block current process until one of its children terminates
 - Return value is the PID of terminated child



system () *Keep your hands off!*

- A naïve approach for our shell program is to just use the system() function
 - Because it takes care of this part for you



pipe

- An interprocess communication channel
 - Implemented as FIFO structures through the kernel
 - Resides in memory

```
:: echo "This is a message" | wc -m  
18
```

The output of **echo** becomes the input for **wc**, data is passed in RAM

- What is another naïve approach to such communication, particularly for our shell program?
 - Use a temporary file **Don't**

pipe

- File descriptor
 - An abstract indicator used to access a file or other I/O resource (pipe, network socket...)
 - A non-negative integer
 - 0 (stdin), 1 (stdout), 2(stderr)

```
[xiaodi_sean]@linux2 ~> (22:48:36 09/26/16)
:: echo "This is a message" 1> msg.txt
Nothing printed
[xiaodi_sean]@linux2 ~> (22:48:43 09/26/16)
:: cat msg.txt
This is a message
```

```
[xiaodi_sean]@linux2 ~> (22:48:45 09/26/16)
:: echo "This is a message" 2> msg.txt
This is a message
[xiaodi_sean]@linux2 ~> (22:50:07 09/26/16)
:: cat msg.txt
Nothing printed
[xiaodi_sean]@linux2 ~> (22:50:09 09/26/16)
```


pipe

- Use `pipe(fd[2])` to create a pipe
 - `fd[0]` refers to the read end of the pipe
 - `fd[1]` the write end
- Terminal commands use `stdin` and `stdout` as I/O by default, we need to “tie” pipes to `stdin/stdout`
 - Use `dup2(old, new)` to make where **new** points to be the copy of where **old** points to, thus they refer to the same I/O port

Background command

- If the parent process doesn't wait, the child process will end up as a zombie after finish: Not good!
- Instead of using the blocking `wait()`, use `waitpid()` with option `WNOHANG` instead.
 - Need to periodically poll the status of the child process
 - Or capture the SIGCHLD signal

Special command

- `cd`
 - Call `chdir()` to do the job.
- `exit`
 - Should we wait for the background process to finish?
 - You can either wait for it or just kill it.

MP4

- Allow users to pipe the standard output from one command to the input of another an arbitrary number of times.
- Open a pipe
 - `pipe()`
- Set stdin and stdout to a pipe
 - `dup2()`
- Reference
 - `man pipe`
 - `man dup2`

MP4

- Allow users to specify whether the process will run in the background or foreground using an '&'. Backgrounding processes should not result in the creation of zombie processes.
 - Cannot use `waitpid(pid, NULL, 0)`, otherwise your shell will not accept new command until the old one finishes.
 - Two options:
 - `waitpid(pid, NULL, WNOHANG)`, a non-blocking version, poll is necessary.
 - Capture the SIGCHLD signal (a parent process will receive this signal whenever its child process exits)

MP4

- Allow your program to take an option "-t" which will run your shell the same way, just without a prompt.
 - No prompt.
- (Bonus) Allow users to specify a custom prompt which supports printing the current directory, username, current date, and current time.
 - A customized prompt.
 - E.g. user:/home/user/svn\$