

CSCE 313

Unix Threads MP7

Wei Zhang

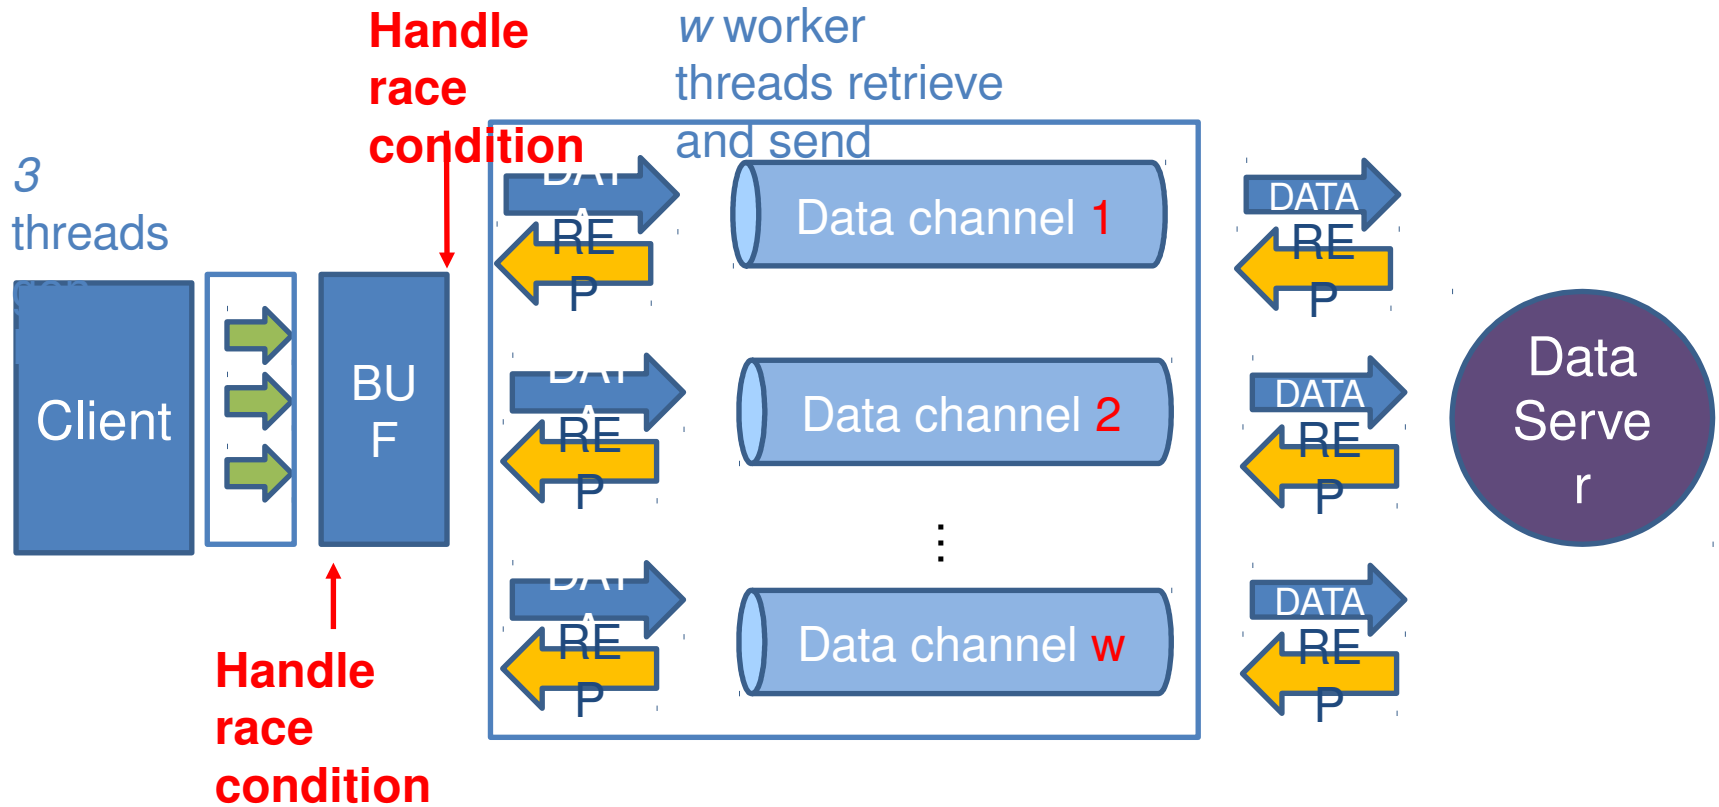
DUE: FRI APR 21, 2017

Acknowledgment: Yi Cui, Wei Zhang, Prof. Tanzir Ahmed, CSCE-3

MP 6: Recap

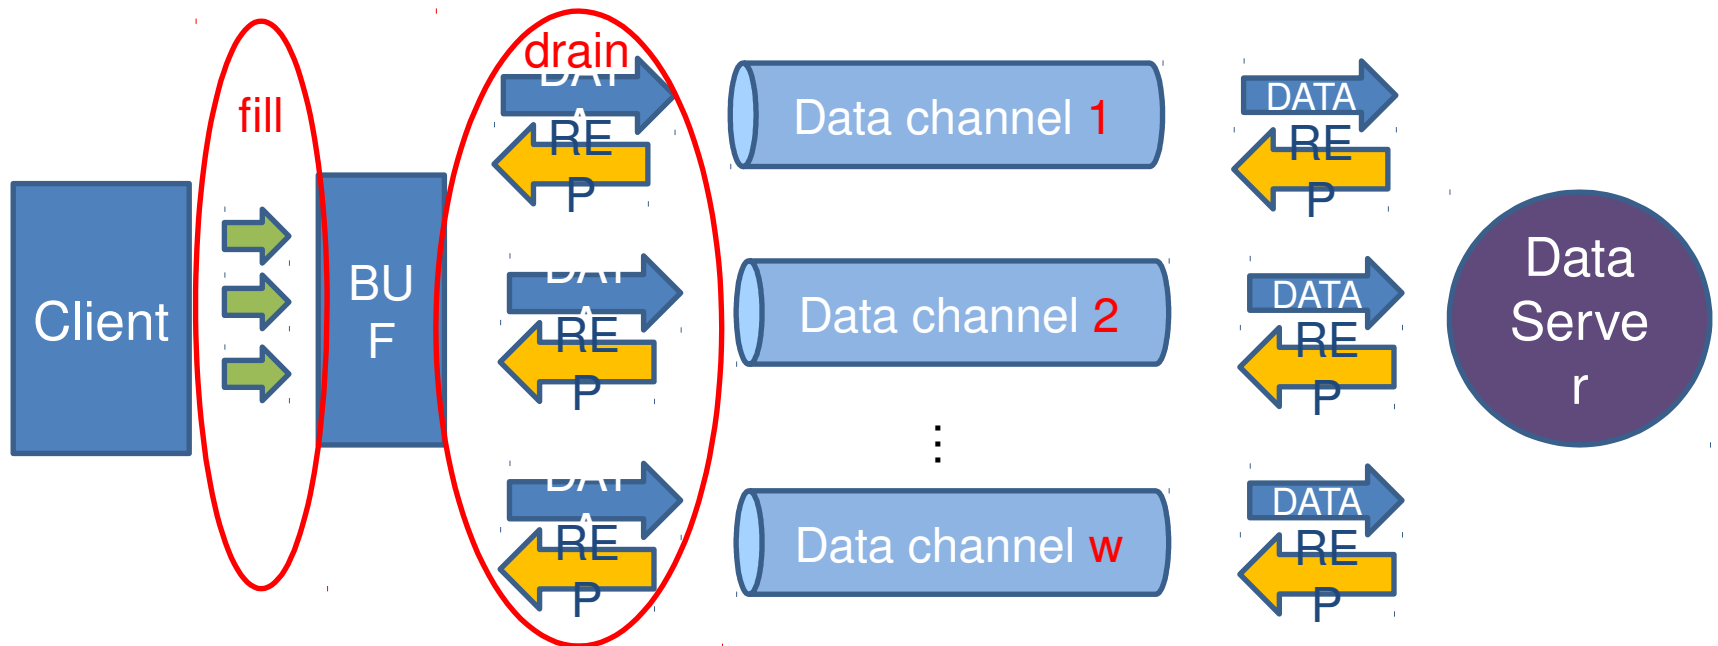
- Client creates 3 threads
 - Generate $3n$ DATA in the buffer, each thread generate one type of data.
 - Run concurrently
 - Feed the data to the buffer, **handle race condition!**
- Then creates w worker threads
 - Each thread creates its own data channel
 - Run concurrently
 - Retrieve data from the buffer
 - Send the data to the server, **handle race condition!**

MP6: recap



MP 6 Limitations

- Limitation #1: Fill the buffer first (with 3 threads), then drain the buffer (using w threads)
 - FIX: Can fill and drain happen simultaneously?



MP 6 Limitations


- Limitation #2: The buffer can be infinite in MP5
 - We have no idea how many data a user can generate
- FIX: Can we limit the size of the buffer?



MP 6 Limitations

- Limitation #3: Poor modularity for worker threads
 - Two distinct roles – (a) interact with server and (b) create stats.

–FIX: Can we devise a scheme to separate these two roles?




NET: What's new in MP7 relative to MP62

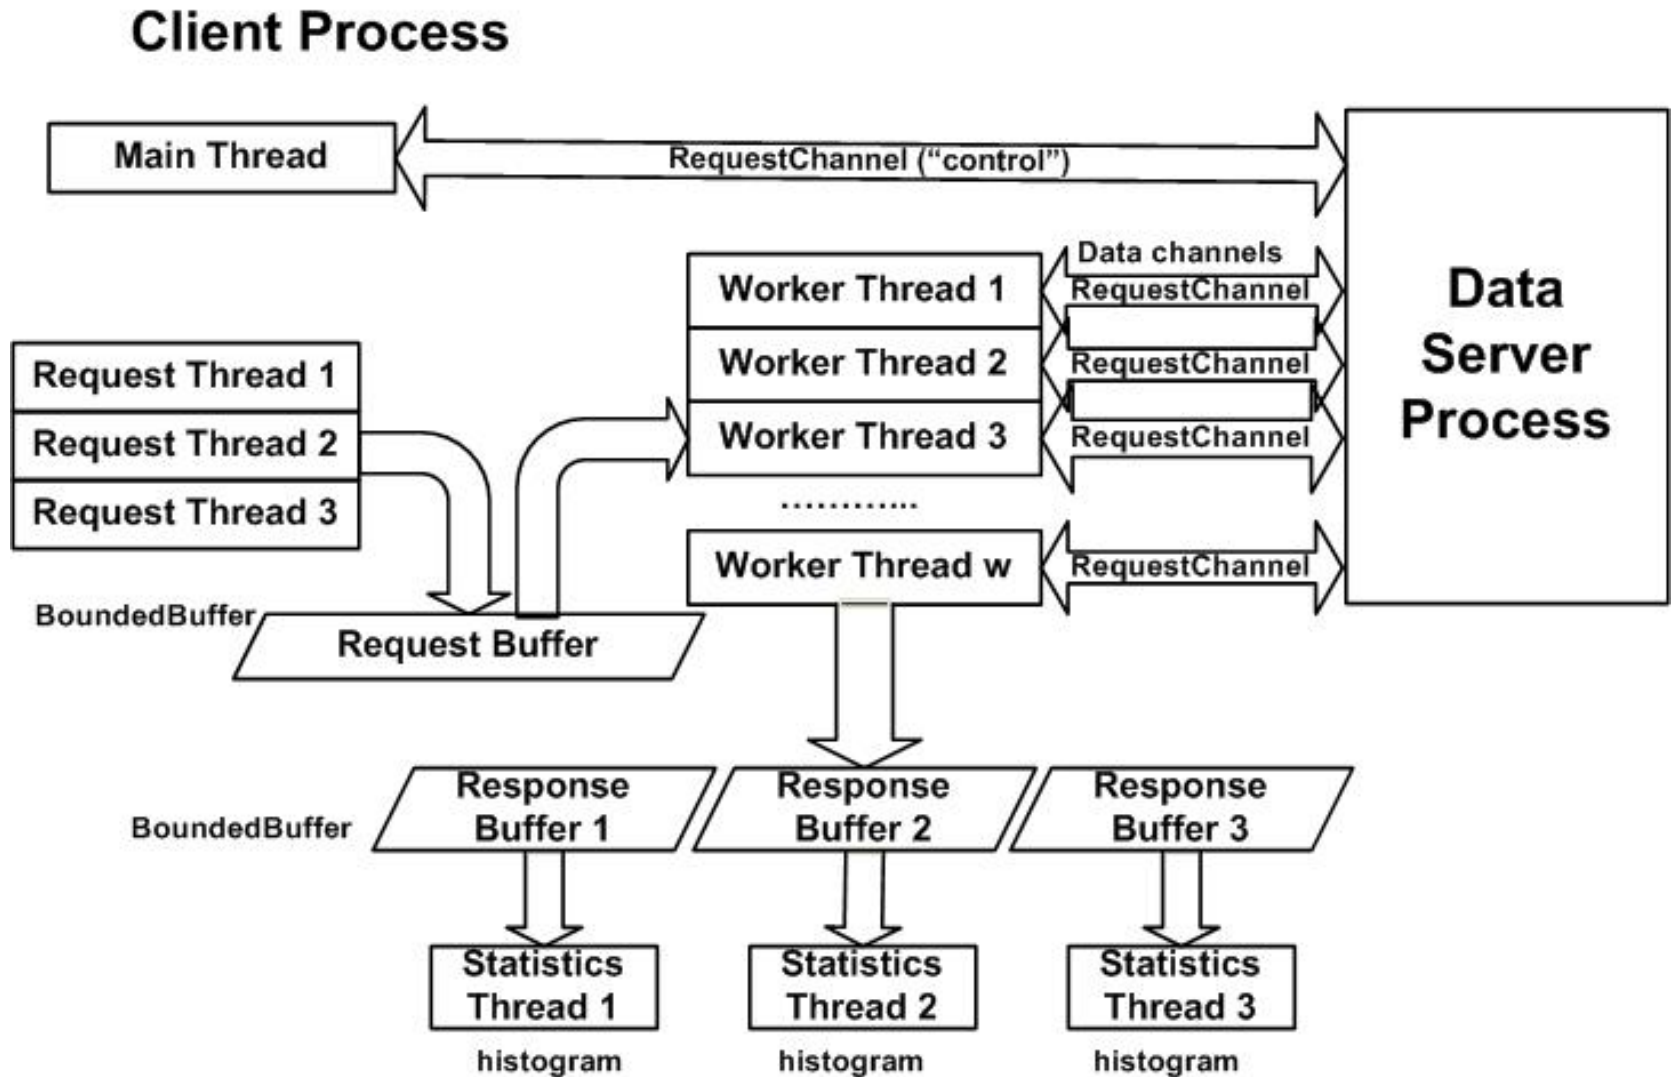
- Request threads can now run concurrently with worker threads
- Request Buffer is BOUNDED
- Worker Threads record server responses in a BOUNDED Response Buffer
- Statistics Threads run concurrently with Worker threads and compile frequency counts from entries in the Response Buffer

MP7 Structure

- Requests from John, Jane, and Joe are mixed into one request buffer
- Workers obtain responses from data server and send them to three separate buffers
- Each person has a dedicated stats thread



Another look at MP7 Structure



MP7 – Tangibles of the Assignment

- `dataserver.cpp`, `reqchannel.h`, `reqchannel.cpp`, and `makefile` are already done for you
- Your task
 - Implement bounded buffer, semaphore classes and client functions in accordance with the requirements published in the “Assignment” section of `MP7_handout.pdf` available [here](#).
 - In addition, write a report with three key sections:
 - Performance Evaluation (especially relative to MP6)
 - Graph the runtime of your client program
 - execution time versus the number of worker threads (i.e. w)
 - execution time versus the size of request buffer (i.e. b)
 - Commentary on your client program performance in context of the system you ran it on

Key Reminders


- Make sure to take note of the specific DON'T(s) mentioned on pp. 4 of MP7_handout.pdf available [here](#).
- Take note of the operations of Semaphore and Bounded Buffer classes mentioned on pp. 4-5 of MP7_handout.pdf available [here](#).
- Consider taking advantage of the BONUS points opportunity mentioned on pp. 5-6 of MP7_handout.pdf available [here](#).
- Read the rubrics carefully! They are listed on the last page pp. 7 of MP7_handout.pdf available [here](#).

USEFUL BACKGROUND MATERIAL



Producer-Consumer

- A classical yet powerful multi-thread programming framework
 - Three parts: producers, consumers, and a bounded buffer
 - Simple and efficient: producers and consumers don't need to talk to each other, they both interact with the queue



Bounded Buffer

- In MP4, we do not set any bounds on queue size, which may lead to two problems
 - Overflow: the queue size can grow to infinity (limited by RAM size)
 - Underflow: consumers may attempt to pop an empty queue
- To solve the problems, we introduce bounded buffers
 - The queue size S is bounded by K . When $S = K$, producers must wait for consumers; when $S = 0$, consumers must wait for producers
 - Use **semaphores** to control the queue size

Semaphore

- You can simply view a semaphore as a **counter**
- When a thread calls `sem.p()`
 - Decrease the counter value C
 - If $C \geq 0$, it can pass; otherwise, it is blocked, i.e., the maximum allowed number has been reached
- When a thread calls `sem.v()`
 - Increase the counter value C
 - If $C \leq 0$, release a blocked thread

Semaphore - Example

- NOT standardized, you need to implement it.
- NOT standardized, you need to implement it.
- Let s be a semaphore
- Let sem be a semaphore
 - Two operations
 - Two operations
 - $sem.p()$
 - $sem.v()$
- A class called semaphore (or other names)
- A class called semaphore (or other names)
 - A mutex variable
 - A int variable $count$
 - A conditional variable
 - A mutex variable mu
 - A conditional variable q

Implementation Example

- *sem.p()*

```
sem.p() mu.lock()
```

```
{  -count --  
  mu.lock();  
  -mutex.unlock()  
  count--;  
  if (count < 0) {  
    cond_wait(&q, &mu)  
  }  
  mu.unlock();  
}
```

Implementation Example

- *sem.p()*

```
sem->mu.lock()
```

```
{  -count --  
  mu.lock();  
  -mutex.unlock()  
  count++;
```

```
  if (count <= 0) { //if Q is empty  
    cond_signal(&q);
```

```
  }
```

```
  mu.unlock();
```

```
}
```

Use semaphore

- Retrieving data (standardized routine)
 - Retrieving data (standardized routine)
 - *full.p()*
 - Remove the data from the buffer
 - *empty.v()*
- Inserting data (standardized routine)
 - Inserting data (standardized routine)
 - *empty.p()*
 - Add the data to the buffer
 - Add the data to the buffer
 - *full.v()*

Producer-Consumer Implementation

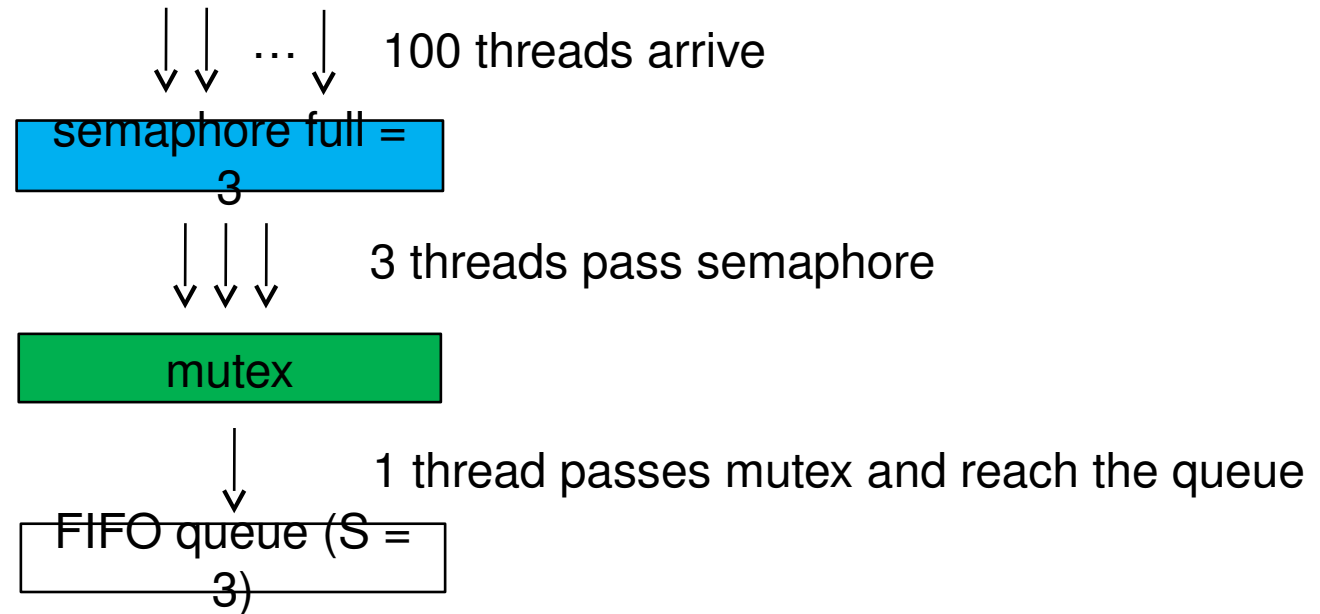
- A standard approach
 - Two semaphores **full** and **empty**, which count the number of items and empty slots in the queue
 - **empty** initialized to buffer capacity, **full** initialized to 0
 - A mutex **m** to control concurrency on queue operations

```
Producer() {  
    while (true) {  
        // decrease the counter of empty slots  
        empty.p();  
        m.lock();  
        add one item to the queue  
        m.unlock();  
        // increase the counter of items  
        full.v();  
    }  
}
```

```
Consumer() {  
    while (true) {  
        // decrease the counter of items  
        full.p();  
        m.lock();  
        pop one item from the queue  
        m.unlock();  
        // increase the counter of empty slots  
        empty.v();  
    }  
}
```

Concurrency Control

- In the consumers perspective, assume there are currently 3 items in the queue (semaphore full = 3)




- Each consumer that reaches the queue must be able to find one item to consume


Parting Reminder ◀◀

- 3 request threads
 - 1 bounded buffer for requests
 - The w worker threads convey requests to the dataserver

Request threads = 3




Bounded buffer



Semaphore to sync

w worker threads



Parting Reminder ◀◀

- 3 statistics threads
 - Draw histogram
 - 3 bounded buffers for response reception
 - The w worker threads put the corresponding responses to the right buffer

