

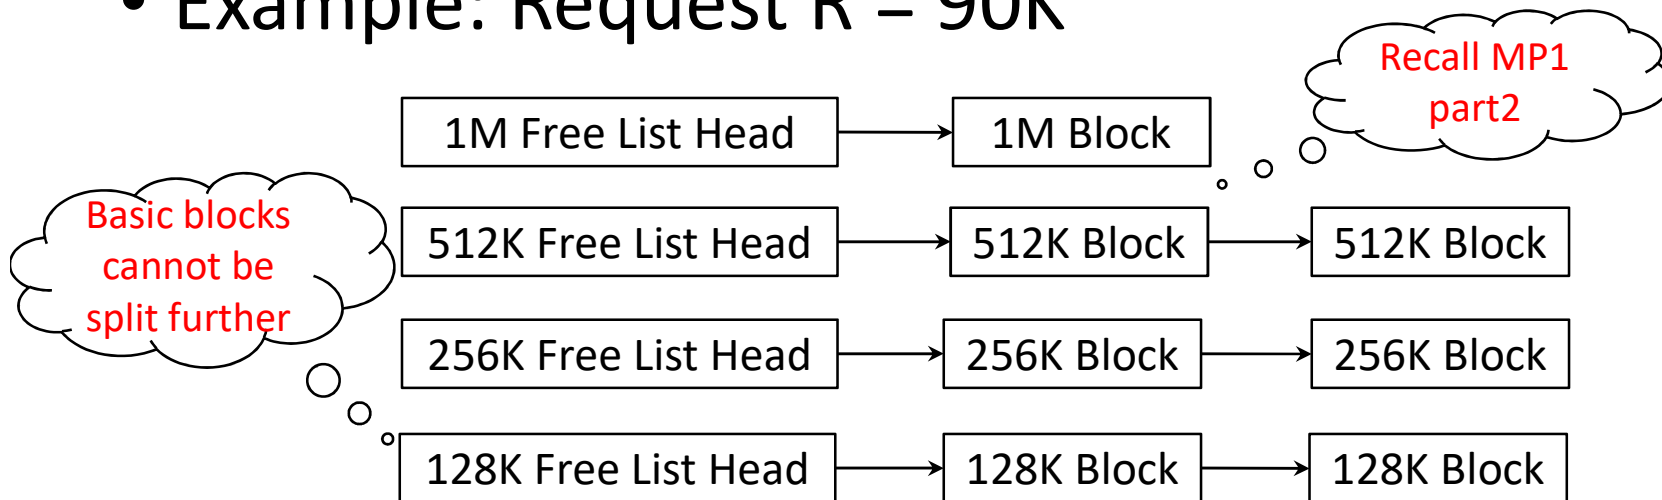


CSCE 313

MPX Introduction


MPX Introduction

- Buddy system is specified by basic block size b and memory size M
 - E.g., $b = 128K$, $M = 1M$
 - Both should be power of 2
- Free blocks are organized into tiered free lists
- Example: Request $R = 90K$



Binary Buddy System

- Simple yet efficient memory management
- Organizes memory in blocks of size 2^k
- When request of size R arrives
 - Find a free block with size that's nearest power of 2
 - If no such block exists, split larger free blocks until a block of correct size is available



- Example
 - 1M memory, request $R_1 = 90K$
 - $R_2 = 150K$, $R_3 = 200K$

Binary Buddy System


- To free a block, check if its buddy is free
 - If so, merge into a larger free block
 - Merge process repeats until we can't go further

- Example

- Free order R_2 , R_1 , R_3

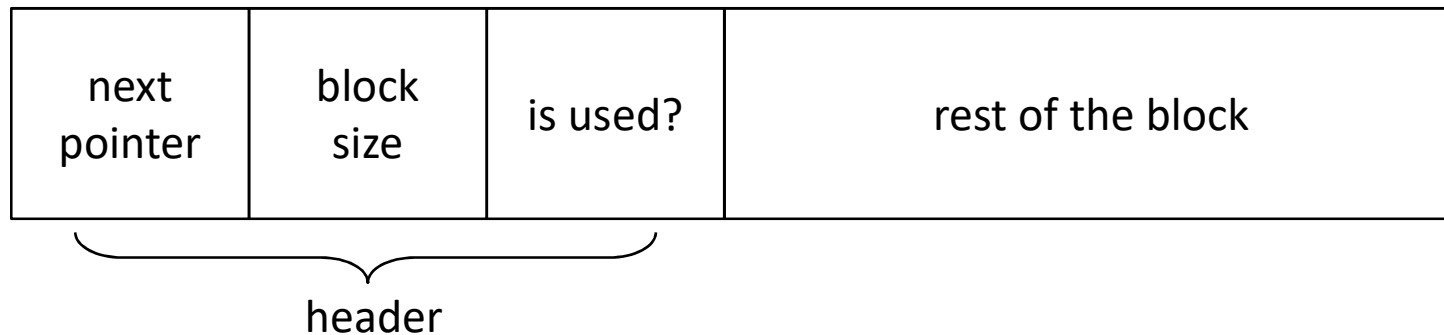
- Drawbacks?

- Internal and external fragmentation R_1
 - Constant splitting and merging



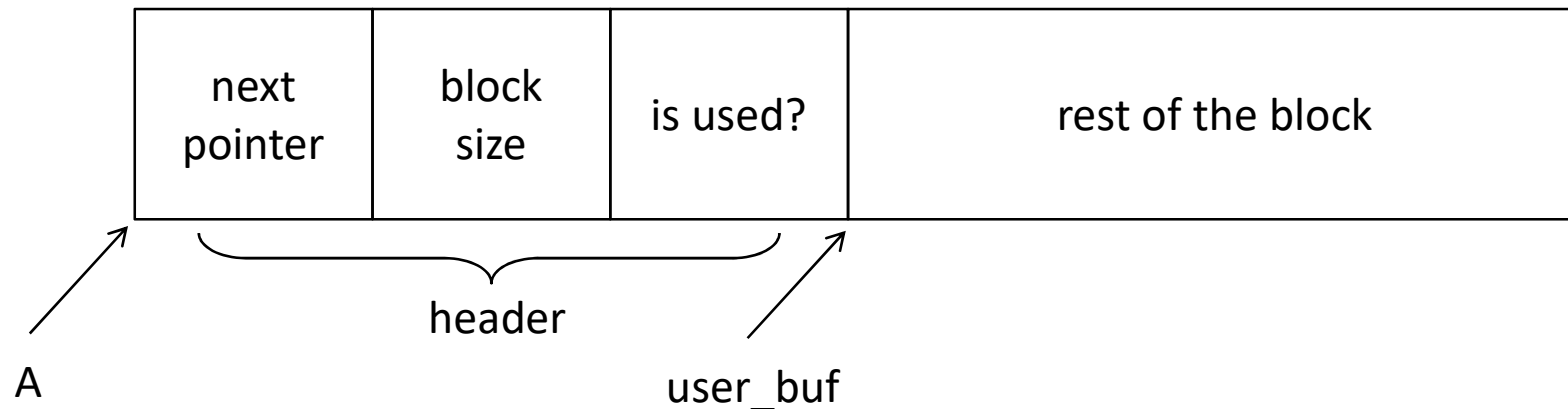
Header

- To organize free blocks into linked lists, we must maintain header information



- Available space for user: $B - \text{sizeof}(\text{header})$
- Example: $R = 120$, $\text{sizeof}(\text{header}) = 16$
 - Size needed: $R + \text{sizeof}(\text{header}) = 136$
 - Round to power of 2: $B = 256$
 - Look for size 256 free blocks


Header




- `void *my_malloc(int R)`
 - Find a free block `A` that satisfies the requested `R`, may need to split larger blocks
 - Setup header info, return pointer `user_buf`
- `bool my_free(void *user_buf)`
 - Locate header `A = user_buf - sizeof(header)`
 - Free this block, may need to merge buddies

Merge Buddies

- Merge two blocks if
 - Both are free
 - They are buddies of the same block size



- Compute buddy addr (for block size $B = 64K$)
 - Approach #1: $A + B = A'$, $A' - B = A$, need extra info in block header to distinguish left and right
 - Approach #2: $A \text{ XOR } B = A'$, $A' \text{ XOR } B = A$, no need to modify header (only works for offset starts from 0)



Wrap up

- `void *my_malloc(int R)`
 - Needed size $B = \text{NearestPow2}(R + \text{sizeof}(\text{header}))$
 - Lookup free list of size B
 - If not empty, allocate the first block, remove it from list
 - If empty, try to find a larger free block and split it all the way down to B
 - Update free lists accordingly
- `bool my_free(void *user_buffer)`
 - Locate header $A = \text{user_buffer} - \text{sizeof}(\text{header})$
 - Compute buddy addr A'
 - $A_{\text{off}} = A - \text{head_pointer}$, $A_{\text{off}}' = A_{\text{off}} \text{ XOR } A \rightarrow \text{block_size}$
 $A' = \text{head_pointer} + A_{\text{off}}'$
 - Merge until we can't go further

init_allocator

```
init_allocator(int b, int M) {  
    if (b > M) error;  
    b = higher_two(b);  
    M = higher_two(M);  
  
    head_pointer = (char*)malloc(M);  allocate memory  
  
    header *h = (header*)head_pointer;  
    h->block_size = M;  
    h->in_use = false;  
    h->next = NULL;  
  
    map[M] = h;  
    for (int i = M / 2; i >= b; i /= 2)  
        map[i] = NULL;  
}
```

} validate input

} setup free lists

my_malloc

```
my_malloc(int R) {  
    int B = higher_two(R + sizeof(header));  
    if (B < basic_block_size) B = basic_block_size;  
    if (B > M) error;  
  
    header *h = NULL; int S = B;  
    while (h == NULL && S <= M) {  
        h = map[S];  
        S *= 2;  
    }  
    if (h == NULL) block not found;  
  
    while (h->block_size > B) split_block(h);  
  
    delete(h);  
    h->in_use = true;  
    return h + 1;  
}
```

compute size needed

find a block that satisfies B

split block until getting size B

delete block from free list
give block to user

split_block

```
split_block(header *h) {  
    delete(h);      delete from free list  
  
    header *h2 = (header*)((char*)h + h->block_size / 2);  
    h->block_size /= 2;  
    h2->block_size = h->block_size;  
    h->in_use = h2->in_use = false;  
  
    add(h);          } add two new blocks to free list  
    add(h2);  
}
```

my_free

```
my_free(Addr a) {  
    header *h = (header*)((char*)a - sizeof(header));  
    h->in_use = false;  
    add(h);  
  
    while (h->block_size < M) {  
        header *buddy = get_buddy(h);  
        if (!buddy->in_use && buddy->block_size == h->block_size)  
            h = combine_blocks(h, buddy);  
        else  
            break; // cannot merge anymore  
    }  
}
```

} add block back to free list

} merge if possible

combine_blocks

```
combine_blocks(header *h1, header *h2) {  
    delete(h1);  
    delete(h2);  
  
    header *h;  
    if (h1 < h2) h = h1;  
    else h = h2;  
    h->block_size *= 2;  
  
    add(h);  
    return h;  
}
```

} delete two small blocks from free list


} merge two blocks into a large block

add large block back to free list

add

```
add(header *node) {  
    int B = node->block_size;  
    node->next = map[B];  
    map[B] = node;  
}
```

} insert node into the beginning of the tier with size B
map[B] points to the head of that tier



delete

```
delete(header *node) {  
    int B = node->block_size;  
    header *head = map[B];  
  
    if (head == node) {  
        map[B] = node->next;  
    }  
    else {  
        while (head->next != node) {  
            head = head->next;  
        }  
        head->next = node->next;  
    }  
}
```

